



Lyftr AI — Backend Assignment

Containerized Webhook API

Objective

Build a production-style FastAPI service that:

- Ingests inbound WhatsApp-like messages exactly once.
 - Validates a simple HMAC-based signature on `/webhook`.
 - Exposes proper liveness/readiness probes.
 - Provides a paginated/filterable `/messages` endpoint.
 - Exposes a Prometheus-style `/metrics` endpoint.
 - Exposes a simple analytical `/stats` endpoint.
 - Uses 12-factor environment configs.
 - Emits structured JSON logs.
 - Runs via Docker Compose using SQLite for storage.
-

Functional Requirements

1. POST `/webhook`

Purpose: Ingest inbound “WhatsApp-like” messages exactly once.

Request:

JSON body:

```
{  
  "message_id": "m1",
```

```
"from": "+919876543210",
"to": "+14155550100",
"ts": "2025-01-15T10:00:00Z",
"text": "Hello"

}
```

- **Headers:**

- **Content-Type:** application/json
- **X-Signature:** <hex HMAC-SHA256 of raw request body using WEBHOOK_SECRET>

Signature:

- **Environment variable:** WEBHOOK_SECRET (non-empty string).

- **Compute:**

```
signature = hex(HMAC_SHA256(secret=WEBHOOK_SECRET,
message=<raw request body bytes>))
```

- **Behavior:**

- If WEBHOOK_SECRET is not set → startup should fail (or /health/ready never becomes 200).
- If X-Signature header is missing or invalid:
 - Return 401 with JSON: {"detail": "invalid signature"}
 - Do not insert anything into DB.
 - Log an error event.
- If X-Signature is valid: proceed to validation + insert.

Validation:

- **message_id:** non-empty string.

- **from / to**: Strings in E.164-like form (start with +, then digits only).
- **ts**: ISO-8601 UTC string with Z suffix, e.g. 2025-01-15T10:00:00Z.
- **text**: optional string, max length 4096.

Invalid payload → 422 (FastAPI/Pydantic style) with a JSON error description, and no DB insert.

Persistence & Idempotency:

- SQLite table **messages** with uniqueness on **message_id**.
- First valid call for a given **message_id**:
 - Insert a row.
 - Return 200 {"status": "ok"}.
- Subsequent calls with the same **message_id** and valid signature:
 - Must not insert a second row.
 - Must still return 200 {"status": "ok"} (idempotent).
 - No stack traces; errors must be handled gracefully.

Response (on success):

```
{
  "status": "ok"
}
```

2. GET /messages

Purpose: List stored messages, with pagination and basic filters.

Query parameters:

- **limit** (optional, int):
 - Default: **50**
 - Min: **1**, Max: **100**
- **offset** (optional, int):
 - Default: **0**
 - Min: **0**
- **from** (optional, string):
 - If present, filter by **from_msisdn** exact match.
- **since** (optional, string):
 - ISO-8601 UTC timestamp.
 - If present, only return messages with **ts >= since**.
- **q** (optional, string):
 - Free-text search in **text** (case-insensitive substring match is enough).

Ordering:

- **Deterministic ordering:** `ORDER BY ts ASC, message_id ASC.`
- **“Newest last” still holds.**

Response shape:

```
{
  "data": [
    {
      "message_id": "m2",
      "from": "+919876543210",
      "to": "+14155550100",
      "text": "Hello"
    }
  ]
}
```

```

        "ts": "2025-01-15T09:00:00Z",
        "text": "Earlier"
    }
    // ...
],
{
    "total": 4,           // total rows matching filters, ignoring
    limit/offset
    "limit": 2,
    "offset": 0
}

```

- **total** must reflect the total number of records for the given filter, not just the number in **data**.
-

3. GET /stats

Purpose: Provide simple message-level analytics.

Behavior: No query params needed for this assignment.

Compute:

```
{
    "total_messages": 123,
    "senders_count": 10,
    "messages_per_sender": [
        { "from": "+919876543210", "count": 50 },
        { "from": "+911234567890", "count": 30 }
        // top senders, up to 10; sorted by count desc
    ],
    "first_message_ts": "2025-01-10T09:00:00Z", // null if no messages
    "last_message_ts": "2025-01-15T10:00:00Z"   // null if no messages
}
```

}

- Implementation can use SQL queries or Python aggregation, but performance should be reasonable for a few thousand rows.
-

4. Health Probes

- **GET /health/live:**
 - Always return 200 once the app is running.
 - **GET /health/ready:**
 - Return 200 only if:
 - DB is reachable and schema is applied.
 - `WEBHOOK_SECRET` is set.
 - Otherwise return 503.
-

5. GET /metrics [OPTIONAL]

Purpose: Expose Prometheus-style metrics.

Format: Text-based Prometheus exposition format.

Minimum required metrics:

- A counter for total HTTP requests, with at least labels for `path` and `status`, e.g.:

```
http_requests_total{path="/webhook", status="200"} 15  
http_requests_total{path="/webhook", status="401"} 2
```

- A counter for webhook processing outcomes, e.g.:

```
webhook_requests_total{result="created"} 10  
webhook_requests_total{result="duplicate"} 5  
webhook_requests_total{result="invalid_signature"} 2
```

```
webhook_requests_total{result="validation_error"} 1
```

- Some latency measurement (bucketed or simple), e.g. buckets:

```
request_latency_ms_bucket{le="100"} 20  
request_latency_ms_bucket{le="500"} 25  
request_latency_ms_bucket{le="+Inf"} 25  
request_latency_ms_count 25
```

- Your exact metric names can differ, but:
 - They must be stable and documented in README.
 - The evaluation script will just check that `/metrics` returns 200 and includes:
 - at least one line starting with `http_requests_total`,
 - and one line starting with `webhook_requests_total`.

6. Structured JSON Logs

- One JSON line per request.
- Required log keys:
 - `ts` (server time, ISO-8601),
 - `level`,
 - `request_id` (unique per request),
 - `method`,
 - `path`,
 - `status`,
 - `latency_ms`.
- For `/webhook` requests, logs must also include:

- `message_id` (when present),
- `dup` (boolean),
- `result("created", "duplicate", "invalid_signature", "validation_error" etc.).`

Logs must be valid JSON per line (good for `jq` / log aggregation).

Non-Functional Requirements

- **Tech:** Python + FastAPI (or similar modern async framework).
- **DB:** SQLite only; DB file must live under a Docker volume (e.g. `/data/app.db`).
- **Validation:** Pydantic (or equivalent). Bad JSON / schema → 422.
- **Idempotency:** Enforced via DB uniqueness (`PRIMARY KEY (message_id)`) plus graceful handling in app layer.
- **Config via env vars:**
 - `DATABASE_URL` (e.g. `sqlite:///data/app.db`).
 - `LOG_LEVEL` (`INFO` / `DEBUG`).
 - `WEBHOOK_SECRET`.
 - No hard-coded paths/secrets.
- **Containerization:**
 - Multi-stage Dockerfile with small runtime image.
 - Docker Compose config with `api` service on `http://localhost:8000`.

Minimal Data Model

```
CREATE TABLE IF NOT EXISTS messages (
    message_id TEXT PRIMARY KEY,
```

```
from_msisdn TEXT NOT NULL,  
to_msisdn  TEXT NOT NULL,  
ts          TEXT NOT NULL,    -- ISO-8601 UTC string  
text        TEXT,  
created_at  TEXT NOT NULL     -- server time ISO-8601  
);  
-- Optional: schema_version table if you want migrations (not  
required for evaluation).
```

Deliverables (Repo Structure)

```
/app  
  main.py           # FastAPI app, middleware, routes  
  models.py         # SQLite init  
  storage.py        # DB operations  
  logging_utils.py # JSON logger  
  metrics.py        # metrics helpers (optional)  
  config.py         # env loading (optional)  
  
/tests  
  test_webhook.py   # valid insert, duplicate, signature cases  
  test_messages.py  # pagination + filters  
  test_stats.py     # stats correctness  
  
Dockerfile  
docker-compose.yml  
Makefile  
README.md
```

Makefile targets:

- make up → docker compose up -d --build

- `make down` → `docker compose down -v`
 - `make logs` → `docker compose logs -f api`
 - `make test` → run your tests (if any)
-

What We Will Actually Run (Evaluation Script Outline)

For each submission, we will:

1. Set env vars and start the stack

```
export WEBHOOK_SECRET="testsecret"  
export DATABASE_URL="sqlite:///data/app.db"
```

```
make up  
# or: docker compose up -d --build  
sleep 10
```

2. Health checks

```
curl -sf http://localhost:8000/health/live >/dev/null  
curl -sf http://localhost:8000/health/ready >/dev/null
```

3. Webhook + Signature (E)

```
BODY=' {"message_id": "m1", "from": "+919876543210", "to": "+14155550100",  
"ts": "2025-01-15T10:00:00Z", "text": "Hello"} '
```

```
# Invalid signature → expect 401  
curl -s -o /dev/null -w "%{http_code}" \  
-H "Content-Type: application/json" \  
-H "X-Signature: 123" \  
-d "$BODY" \  
-X POST http://localhost:8000/webhooks/test
```

```
http://localhost:8000/webhook
# expected: 401

# Compute valid signature
# (We'll use a small helper script – you don't need to provide
this)

Then:

VALID_SIG=<computed from HMAC(SECRET, BODY)>

# Valid signature → 200, row inserted
curl -s -o /dev/null -w "%{http_code}" \
-H "Content-Type: application/json" \
-H "X-Signature: $VALID_SIG" \
-d "$BODY" \
http://localhost:8000/webhook
# expected: 200

# Duplicate with same body + sig → 200, but no new row
curl -s -o /dev/null -w "%{http_code}" \
-H "Content-Type: application/json" \
-H "X-Signature: $VALID_SIG" \
-d "$BODY" \
http://localhost:8000/webhook
# expected: 200
```

4. Seed more messages for `/messages` and `/stats`

We'll send a few more valid messages with different `message_id`, `from`, `ts`, `text` (all with valid signatures).

5. Check /messages pagination & filters

```
# Basic list
curl -s "http://localhost:8000/messages" | jq .

# limit+offset
curl -s "http://localhost:8000/messages?limit=2&offset=0" | jq
'.data | length'

# filter by from=
curl -s "http://localhost:8000/messages?from=+919876543210" | jq .

# filter by since= and q=
curl -s "http://localhost:8000/messages?since=2025-01-15T09:30:00Z"
| jq .
curl -s "http://localhost:8000/messages?q=Hello" | jq .
```

We'll verify that:

- **total** matches the number of rows for that filter.
- **limit/offset** echo back correctly.
- Ordering is **ts** asc, **message_id** asc.

6. Check /stats

```
curl -s "http://localhost:8000/stats" | jq .
```

We'll verify that:

- **total_messages** matches the number of inserted rows.
- **senders_count** is correct.

- `messages_per_sender` entries sum up to `total_messages` (for the senders listed).
- `first_message_ts` and `last_message_ts` are correct min/max.

7. Check `/metrics` [OPTIONAL]

```
curl -s "http://localhost:8000/metrics" | head
```

We'll check:

- HTTP 200.
- Output contains at least:
 - A line starting with `http_requests_total`.
 - A line starting with `webhook_requests_total`.

8. Logs

```
docker compose logs api | head -n 20
```

Quick visual check that:

- Logs are valid JSON per line (we may pipe through `jq`).
- `/webhook` logs include `message_id` and `dup`.

9. Shutdown

```
make down
```

Scoring (10 points)

Most of the scoring will come from the automated checks above.

- Core correctness (4 pts)

- Health endpoints.
 - `/webhook` success + idempotency.
 - Basic `/messages` listing and ordering.
- Advanced endpoints (4 pts)
 - HMAC signature behavior.
 - `/messages` pagination + filters.
 - `/stats` correctness.
- Observability & ops (1 pt)
 - `/metrics` with required metrics. **[OPTIONAL]**
 - JSON logs, `request_id`, `result` fields.
- Docs & hygiene (1 pt)
 - README with:
 - How to run (`make up`, URLs).
 - How to hit endpoints.
 - Brief “Design decisions” section:
 - How you implemented HMAC verification.
 - How your pagination contract works.
 - How you defined `/stats` and metrics.

Constraints & Use of AI

- No external services beyond Docker and Python (no Redis, no Postgres, etc.).
- SQLite only.

- All configuration via environment variables.
- You may use coding assistants (Copilot, ChatGPT, Cursor, etc.).
- In README, add a short note: “Setup Used”
e.g. “VSCode + Copilot + occasional ChatGPT prompts”.

We care about:

- Whether the system actually runs under the above script.
- Whether the behavior matches the exact semantics specified.
- How you structure and explain your solution.

Submission

- Share a GitHub repository link.
- Email to **careers@lyftr.ai** with subject: **Backend Assignment – [Your Name]**